

Project 1B

Compressed Network Communication

INTRODUCTION:

When an application (e.g., shell) is to be used remotely, it is usually not sufficient to simply replace local device I/O (e.g., between the user's terminal and the shell) with network I/O. Remote sessions and network protocols add options and behaviors that did not exist when the application was being used locally. To handle this additional processing without making any changes to the application, it is common to create client-side and server-side agents that handle the network communication and shield the application from the complexities of the remote access protocols. Such intermediary agents can implement valuable features (e.g., encrypting traffic for enhanced security or compressing traffic to improve the performance and reduce the cost of WAN-scale communication), completely transparently to the user and application.

In this project, you will build a multi-process telnet-like client and server. This project is a continuation of Project 1A. It can be broken up into two major steps:

- Passing input and output over a TCP socket
- Compressing communication between the client and server

RELATION TO READING AND LECTURES:

This lab will build on the process and exception discussions, but it is really about researching and exploiting APIs.

PROJECT OBJECTIVES:

- Demonstrate the ability to research new APIs and debug code that exploits them
- Demonstrate ability to do basic network communication
- Demonstrate ability to do exploit a new library
- Develop multi-process debugging skills

DELIVERABLES:

A single tarball (`.tar.gz`) containing:

- C source modules that compile cleanly (with no errors or warnings). You will have (at least) two source files: `lab1b-client.c` and `lab1b-server.c`, which should compile to the executables `lab1b-client` and `lab1b-server`, respectively.
- a `Makefile` to build the program (compiling it with the `-Wall` and `-Wextra` options) and the tarball. It should have distinct targets for the client and server, but it should also have a default target that builds both. It should also have the (now) traditional `clean` and `dist` targets.
- a `README` file describing the contents of the tarball and anything interesting you would like to bring to our attention (e.g., research, limitations, features, testing methodology).

PROJECT DESCRIPTION:

Study the following documentation:

- [socket\(7\)](#) ... for creating communication endpoints (networking)
- [zlib\(3\)](#) and the [zlib tutorial](#) ... a library for performing [gzip](#) compression and decompression (based on a combination of [LZW](#) and [Huffman](#) coding)

Passing input and output over a TCP socket

Using Project 1A's `--shell` option as a starting point, create a client program (`lab1b-client`) and a server program (`lab1b-server`), both of which support a mandatory `--port=port#` switch. (Note: you can remove the `--shell` option)

Your overall program design will look like the diagram above. You will have a client process, a server process, and a shell process. The first two are connected via TCP connection, and the last two are connected via pipes.

The client program

- The client program will open a connection to a server (port specified with the mandatory `--port=` command line parameter) rather than sending it directly to a shell. The client should then send input from the keyboard to the socket (while echoing to the display), and input from the socket to the display.
- Maintain the same non-canonical (character at a time, no echo) terminal behavior you used in Project 1A.
- Include, in the client, a `--log=filename` option, which maintains a record of data sent over the socket (see below).
- If a `^D` is entered on the terminal, simply pass it through to the server like any other character.
- If a `^C` is entered on the terminal, simply pass it through to the server like any other character.
- In the second part of this project we will add a `--compress` switch to the client.
- If the client encounters an unrecognized argument it should print an informative usage message (on `stderr`) and exit with a return code of 1. If any system call fails, it should print an informative error message (on `stderr`) and exit with a return code of 1.
- Before exiting, restore normal terminal modes.

For all of this testing, the client and server will be running on the same host (`localhost`), but if you would like to add a `--host=name` parameter you should be able to access a shell session from other computers.

The server program

- The server program will connect with the client, receive the client's commands and send them to the shell, and will "serve" the client the outputs of those commands.
- The server program will listen on a network socket (port specified with the mandatory `--port=` command line parameter).
- Accept a connection when it is made.
- Once a connection is made, fork a child process, which will exec a shell to process commands received; the server process should communicate with the shell via pipes (as you did in Project 1A).
- Redirect the shell process's `stdin/stdout/stderr` to the appropriate pipe ends (similar to what you did in Project 1A).
- Input received through the network socket should be forwarded through the pipe to the shell. Because the server forks the shell (and knows its process ID), the processing of `^C` (turning it into a `SIGINT` to the shell) must be done in the server. Similarly, when the server encounters a `^D` it should close the write side of the pipe to the shell (as was done in project 1A).
- Input received from the shell pipes (which receive both `stdout` and `stderr` from the shell) should be forwarded out to the network socket.
- If the server gets an EOF or `SIGPIPE` from the shell pipes (e.g., the shell exits), harvest the shell's completion status, log it to `stderr` (as you did in project 1A), and exit with a return code of 0.
- If the server gets an EOF or read error from the network connection, (e.g., the client exits), close the write pipe to the shell, await an EOF from the pipe from the shell, and harvest and report its termination status (as

above).

- If the server encounters an unrecognized argument it should print an informative usage message (on stderr) and exit with return code of 1. If any system call fails, it should print an informative error message (on stderr) and exit with an return code of 1.
- In the second part of this project we will add a **--compress** switch to the server.

To send commands to the shell, you will need to start both the client and the server. On the client side, the behavior should be very similar to that of Project 1A, in that you will see your commands echoed back to the terminal as well as the output from the shell. Do not display anything (other than error messages) on the server side.

If you are not familiar with socket programming, you may find this [tutorial](#) to be useful. (Note that we are not expecting you to become expert in how compression in general or this particular form of encryption works. Your task is to understand it well enough to use it.)

Choose a random port number (greater than 1024) for your testing. Port numbers below 1024 are reserved. Be advised that listeners may remain active for several seconds after the server exits. This means that you may get occasional socket-in-use errors if you restart the server before the old listen is taken down.

Because the server is started independently from the client, it will have its own stderr and stdout which go the terminal session from which it was started. If you want to implement debug output from the server, you can send it to the server's stdout (file descriptor 1).

Shut-down Order

If the client initiates the closing of the socket, it may not receive the last output from the shell. To ensure that no output is lost, shut-downs should be initiated from the server side:

1. an *exit(1)* command is sent to the shell, or the server closes the write pipe to the shell.
2. the shell exits, causing the server to receive an EOF on the read pipe from the shell.
3. the server collects and reports the shell's termination status.
4. the server closes the network socket to the client, and exits.
5. the client continues to process output from the server until it receives an error on the network socket from the server.
6. the client restores terminal modes and exits.

After reporting the shell's termination status and closing the network socket, the server should exit. Otherwise it would tie up the socket and prevent testing new server versions. After your test is complete, the client, server, and shell should all be gone. There should be no remaining orphan processes.

The --log option

To ensure that compression is being correctly done, we will ask you to add a **--log=filename** option to your client. If this option is specified, all data written to or read from the server should be logged to the specified file. Prefix each log entry with **SENT # bytes:** or **RECEIVED # bytes:** as appropriate. (Note the colon and space between the word **bytes** and the start of the data). Each of these lines should be followed by a newline character (even if the last character of the reported string was a newline).

Before running your program in a mode that creates a log file, please use the *ulimit* command to ensure that your log file does not get too large, which could happen if you have a bug in your program. *ulimit 10000* should be sufficient, but check out the [ulimit man page](#) for further details. Failing to limit the size of log files has filled up the HSSEAS Linux servers' /tmp directories in the past, which makes the machines hard to use for everybody, including you.

Sample log format output:

SENT 35 bytes: sendingsendingsendingsendingsending
RECEIVED 18 bytes: receivingreceiving

Compressed communication

The purpose of compression is to reduce the amount of space data takes up. In communications scenarios, a compressed version of data will use less bandwidth to transmit the same data than its uncompressed form would use. On the other hand, compression will require processing at the sending and receiving end before the data can be effectively used at the receiving process. Sometimes the benefits in reduced bandwidth are worth the costs of extra processing and sometimes they are not, so compression is not used by default in network communications. While compression makes data difficult to read, unlike encryption it does not provide any privacy for the data, since anyone who obtains the compressed version of the data can use standard algorithms to decompress it. Thus, compression is used for efficiency, not for privacy. When both efficiency and privacy are desired, data is first compressed, then encrypted.

In this project, you will only perform compression, without encryption. You will use a standard compression library to improve efficiency of your socket communications.

- Depending on what system you do your development on you may need to install the `zlib` package to get the compression library and its documentation.
- When working with a new library, it is often easiest to start with a very simple application. A multi-process application like our client, server, and shell is intrinsically complex. You may find it much easier to get the compression/decompression code working in a simple program that reads from and writes to a pipe. Then, after you have it working, you can integrate that code into your client and server.
- Add a **--compress** command line option to your client and server which, if included, will enable compression (of all traffic in both directions).
- Modify both the client and server applications to compress traffic before sending it over the network and decompress it after receiving it.
Note that you are not sending a continuous data stream in either direction. Rather you are sending a series of independent messages (each of which must be complete and an integer number of bytes long). Make sure that you choose the appropriate `z_SYNC` option when you are doing your compression (and decompression).
- Using these options, run a session and use the **--log** option to verify that compression and decompression are working properly. The **--log** option should record outgoing data post-compression and incoming data pre-decompression. You should expect to see:
 1. no clear text going either to or from the shell
 2. a reduction in the number of bytes sent in both directions
- Do not include your **--log** files with your submission; we will generate our own from your program.

The compression algorithm you are using is a progressive one, which becomes more effective as it processes more data. As a result you may note very little reduction in the character-at-a-time commands being sent to the shell. Larger output coming back from the shell should, however, compress quite effectively.

Summary of exit codes

0 ... normal execution, shutdown on ^D
1 ... unrecognized argument or system call failure

SUBMISSION:

Your **README** file must include lines of the form:

NAME: *your name*
EMAIL: *your email*
ID: *your student ID*

And, if slip days are allowed on this project, and you want to use some, this too must be included in the **README** file:

SLIPDAYS: *#days*

If, for instance, you wanted to use two slip-days, you would add the following line:

SLIPDAYS: 2

Your name, student ID, and email address should also appear as comments at the top of your `Makefile` and each source file.

Your tarball should have a name of the form `lab1b-studentID.tar.gz`. You can sanity check your submission with this [test script](#).

We will test it on a departmental Linux server. You would be well advised to test your submission on that platform before submitting it.

GRADING:

Points for this project will be awarded:

value feature

Packaging and build (15% total)

- 5% un-tars expected contents
- 5% clean build of correct programs w/default action (no warnings)
- 3% Makefile has working `clean` and `dist` targets
- 2% reasonableness of `README` contents

Communication (50% total)

- 5% input properly echos (character at a time) from client
- 5% client passes (character at a time) data between server and terminal
- 5% server passes (character at a time) data between client and shell
- 5% proper translation of input CR (from keyboard) into NL for shell input
- 5% proper translation of NL (from) into CR NL to display
- 5% `^D` from client closes connection, restores terminal modes, exits `RC=0`
- 5% `^C` from client sends a `SIGINT` to shell
- 5% server properly logs shell termination status
- 5% `--log=` option properly reports sent data
- 5% `--log` option properly reports received data

Compression (25% total)

- 5% `--compress` reduces number of bytes sent over the wire in both directions
- 5% data is compressed before sending (tested with `--log` option)
- 5% data is decompressed after receiving (tested with `--log` option)
- 5% compressed commands sent to shell are properly decompressed before execution
- 5% compressed output sent by shell are properly decompressed before display

Code Review (10%)

- 10% correct library use